

```
#Importing required libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from scipy import stats as st
```

```
# Reading the .xlsx file
```

```
df=pd.read_csv(r"C:\Users\saidu\Downloads\data.xlsx - Sheet1.csv")
df.drop("Unnamed: 0",axis=1,inplace=True)
df.head()
```

	ID	Salary	DOJ	DOL
Designation \				
0	203097	420000	2012-06-01	present senior quality engineer
1	579905	500000	2013-09-01	present assistant manager
2	810601	325000	2014-06-01	present systems engineer
3	267447	1100000	2011-07-01	present senior software engineer
4	343523	200000	2014-03-01	2015-03-01 00:00:00

	JobCity	Gender	DOB	10percentage
10board \				
0	Bangalore	f	1990-02-19	84.3 board ofsecondary education,ap
1	Indore	m	1989-10-04	85.4
2	Chennai	f	1992-08-03	85.0
3	Gurgaon	m	1989-12-05	85.6
4	Manesar	m	1991-02-27	78.0

	ComputerScience	MechanicalEngg	ElectricalEngg	TelecomEngg
0	-1	-1	-1	-1
1	-1	-1	-1	-1
2	-1	-1	-1	-1
3	-1	-1	-1	-1
4	-1	-1	-1	-1

	CivilEngg	conscientiousness	agreeableness	extraversion
0	-1	0.9737	0.8128	0.5269
1.35490				
1	-1	-0.7335	0.3789	1.2396

```

0.10760
2      -1      0.2718      1.7109      0.1637      -
0.86820
3      -1      0.0464      0.3448      -0.3440      -
0.40780
4      -1      -0.8810      -0.2793      -1.0697
0.09163

```

```

    openness_to_experience
0      -0.4455
1      0.8637
2      0.6721
3     -0.9194
4     -0.1295

```

```
[5 rows x 38 columns]
```

#Shape of the given data

```
df.shape
```

```
(3998, 38)
```

Information about the AMCAT data

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 3998 entries, 0 to 3997
```

```
Data columns (total 38 columns):
```

#	Column	Non-Null Count	Dtype
0	ID	3998 non-null	int64
1	Salary	3998 non-null	int64
2	DOJ	3998 non-null	datetime64[ns]
3	DOL	3998 non-null	object
4	Designation	3998 non-null	object
5	JobCity	3998 non-null	object
6	Gender	3998 non-null	object
7	DOB	3998 non-null	datetime64[ns]
8	10percentage	3998 non-null	float64
9	10board	3998 non-null	object
10	12graduation	3998 non-null	int64
11	12percentage	3998 non-null	float64
12	12board	3998 non-null	object
13	CollegeID	3998 non-null	int64
14	CollegeTier	3998 non-null	int64
15	Degree	3998 non-null	object
16	Specialization	3998 non-null	object
17	collegeGPA	3998 non-null	float64
18	CollegeCityID	3998 non-null	int64
19	CollegeCityTier	3998 non-null	int64

20	CollegeState	3998	non-null	object
21	GraduationYear	3998	non-null	int64
22	English	3998	non-null	int64
23	Logical	3998	non-null	int64
24	Quant	3998	non-null	int64
25	Domain	3998	non-null	float64
26	ComputerProgramming	3998	non-null	int64
27	ElectronicsAndSemicon	3998	non-null	int64
28	ComputerScience	3998	non-null	int64
29	MechanicalEngg	3998	non-null	int64
30	ElectricalEngg	3998	non-null	int64
31	TelecomEngg	3998	non-null	int64
32	CivilEngg	3998	non-null	int64
33	conscientiousness	3998	non-null	float64
34	agreeableness	3998	non-null	float64
35	extraversion	3998	non-null	float64
36	nueroticism	3998	non-null	float64
37	openess_to_experience	3998	non-null	float64

dtypes: datetime64[ns](2), float64(9), int64(18), object(9)
memory usage: 1.2+ MB

Exploratory Data Analysis

Getting the insights from the data which includes

- Missing values
- Duplicated values
- Ouliers
- Distributions
- Relationships

```
# Chacking missing values
df.isna().sum()
```

ID	0
Salary	0
DOJ	0
DOL	0
Designation	0
JobCity	0
Gender	0
DOB	0
10percentage	0
10board	0
12graduation	0
12percentage	0
12board	0

```

CollegeID      0
CollegeTier    0
Degree         0
Specialization 0
collegeGPA     0
CollegeCityID  0
CollegeCityTier 0
CollegeState   0
GraduationYear 0
English        0
Logical        0
Quant          0
Domain         0
ComputerProgramming 0
ElectronicsAndSemicon 0
ComputerScience 0
MechanicalEngg 0
ElectricalEngg 0
TelecomEngg    0
CivilEngg      0
conscientiousness 0
agreeableness  0
extraversion   0
nueroticism    0
openess_to_experience 0
dtype: int64

# Chacking duplicated values
df.duplicated().sum()

0

```

Univariate Analysis

- Analysing the data using single feature/variable.

```

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3998 entries, 0 to 3997
Data columns (total 38 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   ID                    3998 non-null   int64
 1   Salary                3998 non-null   int64
 2   DOJ                  3998 non-null   datetime64[ns]
 3   DOL                  3998 non-null   object
 4   Designation           3998 non-null   object
 5   JobCity               3998 non-null   object

```

6	Gender	3998	non-null	object
7	DOB	3998	non-null	datetime64[ns]
8	10percentage	3998	non-null	float64
9	10board	3998	non-null	object
10	12graduation	3998	non-null	int64
11	12percentage	3998	non-null	float64
12	12board	3998	non-null	object
13	CollegeID	3998	non-null	int64
14	CollegeTier	3998	non-null	int64
15	Degree	3998	non-null	object
16	Specialization	3998	non-null	object
17	collegeGPA	3998	non-null	float64
18	CollegeCityID	3998	non-null	int64
19	CollegeCityTier	3998	non-null	int64
20	CollegeState	3998	non-null	object
21	GraduationYear	3998	non-null	int64
22	English	3998	non-null	int64
23	Logical	3998	non-null	int64
24	Quant	3998	non-null	int64
25	Domain	3998	non-null	float64
26	ComputerProgramming	3998	non-null	int64
27	ElectronicsAndSemicon	3998	non-null	int64
28	ComputerScience	3998	non-null	int64
29	MechanicalEngg	3998	non-null	int64
30	ElectricalEngg	3998	non-null	int64
31	TelecomEngg	3998	non-null	int64
32	CivilEngg	3998	non-null	int64
33	conscientiousness	3998	non-null	float64
34	agreeableness	3998	non-null	float64
35	extraversion	3998	non-null	float64
36	nueroticism	3998	non-null	float64
37	openess_to_experience	3998	non-null	float64


dtypes: datetime64[ns](2), float64(9), int64(18), object(9)
memory usage: 1.2+ MB

What is the distribution of Salary

```
pd.DataFrame(df["Salary"].describe())
```

	Salary
count	3.998000e+03
mean	3.076998e+05
std	2.127375e+05
min	3.500000e+04
25%	1.800000e+05
50%	3.000000e+05
75%	3.700000e+05
max	4.000000e+06

```
sns.kdeplot(data=df["Salary"])
plt.grid()
plt.title("Distribution of Salary")
plt.show()
```



Insights

- In between 0 to 100000 the salaries are more compared to other salaries.
- After 300000 there are less salaries.

What is the average collegeGPA of students?

```
df["collegeGPA"].mean()
71.48617058529265
```

What are the counts of different JobCity values?

```
pd.DataFrame(df["JobCity"].value_counts())
```

	count
JobCity	
Bangalore	627
-1	461
Noida	368
Hyderabad	335
Pune	290
...	...
Tirunelveli	1
Ernakulam	1
Nanded	1
Dharmapuri	1
Asifabadbangalore	1

[339 rows x 1 columns]

Which Specialization is most common among the students?

```
df["Specialization"].value_counts().head(10)
```

Specialization	
electronics and communication engineering	880
computer science & engineering	744
information technology	660
computer engineering	600
computer application	244
mechanical engineering	201
electronics and electrical engineering	196
electronics & telecommunications	121
electrical engineering	82
electronics & instrumentation eng	32

Name: count, dtype: int64

```
d1=pd.DataFrame(df["Specialization"].value_counts().head(10))
d1
```


	count
Specialization	
electronics and communication engineering	880
computer science & engineering	744
information technology	660
computer engineering	600
computer application	244
mechanical engineering	201
electronics and electrical engineering	196
electronics & telecommunications	121

```

electrical engineering      82
electronics & instrumentation eng  32

sns.barplot(y=d1.index,x=d1["count"],hue=d1.index)
plt.title("Distribution of Specilization")
plt.show()

```



Insights

- There are more electronics engineers compared to others.
- There are less electrical,instrumentation engineers..

```

# Assuming df is your DataFrame
# Set up the number of subplots based on the number of columns
n_cols = len(df.columns)
n_rows = int(np.ceil(n_cols / 3)) # 3 columns per row for better layout

fig, axes = plt.subplots(n_rows, 3, figsize=(20, n_rows * 6))
axes = axes.flatten() # Flatten the axes array for easier indexing

# Iterate over each column in the DataFrame and each subplot axis
for i, col in enumerate(df.columns):

    # Check if the column is categorical
    if df[col].dtype == 'object' or df[col].dtype.name == 'category':
        # Categorical column - use countplot
        sns.countplot(x=col, data=df, ax=axes[i])
        axes[i].set_title(f'Distribution of {col} (Categorical)')

```



```

# Check if the column is datetime
elif pd.api.types.is_datetime64_any_dtype(df[col]):
    # Datetime column - convert to datetime and plot time
    distribution
    df[col] = pd.to_datetime(df[col])
    df[col].value_counts().sort_index().plot(ax=axes[i])
    axes[i].set_title(f'Time Distribution of {col} (Datetime)')
    axes[i].set_xlabel('Date')
    axes[i].set_ylabel('Count')

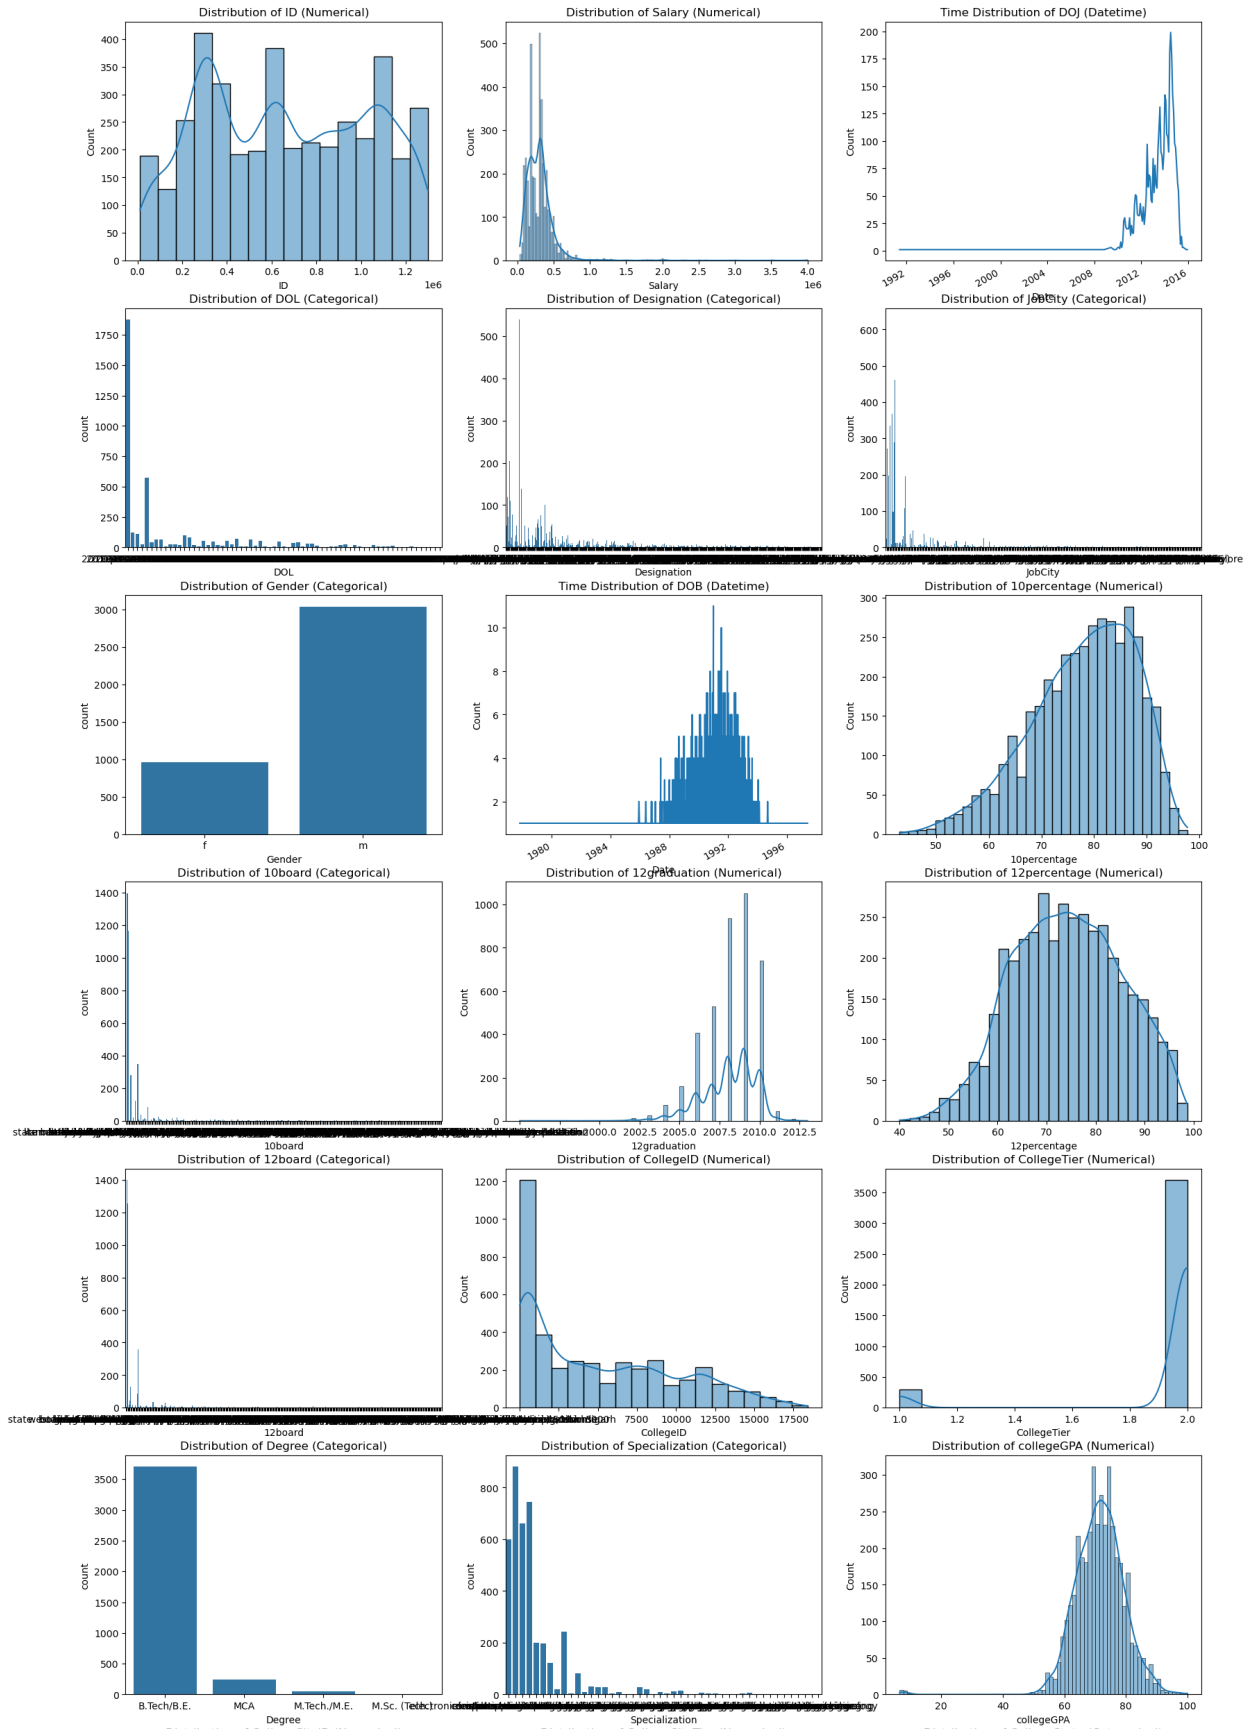
# Check if the column is numerical
elif pd.api.types.is_numeric_dtype(df[col]):
    # Numerical column - use histplot
    sns.histplot(df[col], kde=True, ax=axes[i])
    axes[i].set_title(f'Distribution of {col} (Numerical)')

# Hide unused axes if fewer columns than subplots
if i >= n_cols:
    axes[i].axis('off')


# Adjust layout for better spacing between subplots

plt.show()


```




```
for i in df.columns:  
    if df[i].dtype=="int" or df[i].dtype=="float":  
        sns.boxplot(x=df[i])  
        plt.title("Boxplot for {}".format(i))  
        plt.show()
```




Boxplot for 12percentage




Boxplot for collegeGPA




Boxplot for Domain




Boxplot for conscientiousness




Boxplot for agreeableness

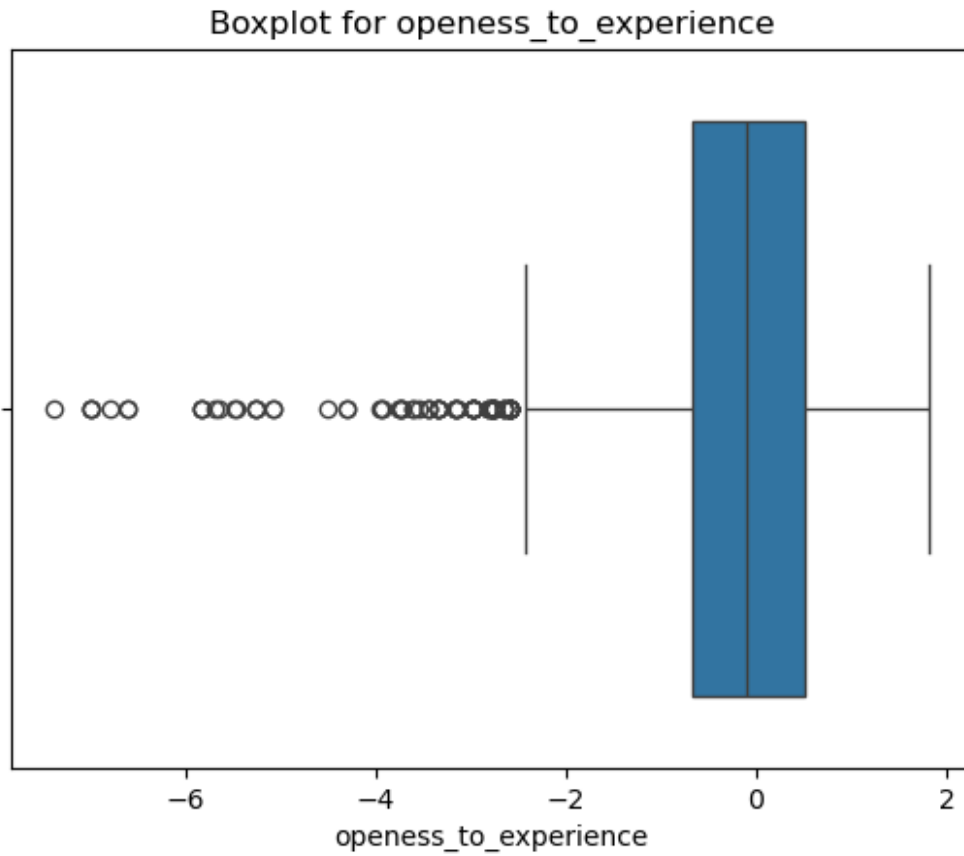


Boxplot for extraversion




Boxplot for nueroticism






```
for i in df.columns:
    if df[i].dtype=="object":
        sns.barplot(x=df[i].unique()[:10],y=df[i].value_counts()[:10])
        plt.title("Distribution of {}".format(i))
        plt.xticks(rotation=90)
        plt.show()
```


Distribution of DOL





Distribution of Designation




Distribution of JobCity










Distribution of 12board









Bivariate Analysis

- Analysing the data using two features.

How does collegeGPA vary across different Specialization?

```
g1=df.groupby("Specialization")  
[["collegeGPA"]].mean().sort_values(by="collegeGPA",ascending=False)  
g1
```


Specialization	collegeGPA
embedded systems technology	88.000000
control and instrumentation engineering	82.100000
information science	81.200000
internal combustion engine	80.600000
industrial & management engineering	80.000000
computer science	77.385000

computer and communication engineering	77.260000
power systems and automation	76.000000
other	75.619231
metallurgical engineering	75.550000
information & communication technology	75.500000
instrumentation and control engineering	75.380000
telecommunication engineering	74.776667
mechatronics	74.375000
industrial engineering	73.850000
computer application	73.700779
mechanical and automation	73.530000
biotechnology	73.155333
industrial & production engineering	73.146000
electrical engineering	72.820000
polymer technology	72.790000
civil engineering	72.761034
automobile/automotive engineering	72.690000
electronics & instrumentation eng	72.679063
electronics and communication engineering	72.126170
electronics and electrical engineering	72.097143
ceramic engineering	72.000000
applied electronics and instrumentation	71.888889
computer science & engineering	71.779798
electronics and instrumentation engineering	71.634815
computer engineering	71.046500
electronics	71.000000
information technology	70.510803
chemical engineering	70.138889
computer networking	70.130000
mechanical engineering	70.109154
computer science and technology	69.091667
electronics & telecommunications	69.020413
aeronautical engineering	68.033333
instrumentation engineering	67.547500
information science engineering	67.322593
electronics and computer engineering	67.313333
biomedical engineering	64.650000
electronics engineering	61.318947
mechanical & production engineering	58.000000
electrical and power engineering	35.705000

```

sns.barplot(y=g1.index[:10],x=g1["collegeGPA"][:10],hue=g1.index[:10])
plt.title("Distribution of CollgGPA in Specialization")
plt.show()

```




Insights

- The Average GPA of embedded systems is more compared to others
- There are less GPA for others, metallurgical engineering compared to others..

Relationship between Gender and Salary?

```
sns.boxplot(y=df["Gender"], x=df["Salary"], hue=df["Gender"])  
plt.show()
```



Does the GraduationYear impact JobCity selection?

```
g2=pd.crosstab(index=df["GraduationYear"],columns=df["JobCity"],margin
s=True,margins_name="Total")
```

```
g2
```

JobCity \ GraduationYear	-1	Chennai	Delhi	Mumbai	Pune	ariyalur
0	0	0	0	0	0	0
2007	0	0	0	0	0	0
2009	1	0	0	0	0	0
2010	16	0	0	1	0	1
2011	44	0	0	0	0	0
2012	115	1	0	0	1	0
2013	170	0	1	1	0	0
2014	108	0	0	0	0	0

0						
2015	6	0	0	0	0	0
0						
2016	0	0	0	0	0	0
0						
2017	1	0	0	0	0	0
0						
Total	461	1	1	2	1	1
1						

JobCity mumbai A-64,sec-64,noida AM ... shahibabad
singaruli \

GraduationYear					...	
0	0		0	0	...	0
0						
2007	0		0	0	...	0
0						
2009	0		0	0	...	0
1						
2010	0		0	0	...	0
0						
2011	0		0	0	...	0
0						
2012	0		0	0	...	0
0						
2013	1		0	0	...	1
0						
2014	0		1	1	...	0
0						
2015	0		0	0	...	0
0						
2016	0		0	0	...	0
0						
2017	0		0	0	...	0
0						
Total	1		1	1	...	1
1						

JobCity	sonepat	thane	trivandrum	udaipur	vapi	vizag	\
GraduationYear							
0	0	0	0	0	0	0	
2007	0	0	0	0	0	0	
2009	0	0	0	0	0	0	
2010	1	0	0	0	0	0	
2011	0	1	0	1	0	0	
2012	0	0	0	1	0	0	
2013	0	0	1	0	0	0	
2014	0	0	1	0	1	1	
2015	0	0	0	0	0	0	

2016	0	0	0	0	0	0
2017	0	0	0	0	0	0
Total	1	1	2	2	1	1

JobCity	vsakhapttnam	Total
GraduationYear		
0	0	1
2007	0	1
2009	0	24
2010	1	292
2011	0	507
2012	0	847
2013	0	1181
2014	0	1036
2015	0	94
2016	0	7
2017	0	8
Total	1	3998

[12 rows x 340 columns]


Does Designation affect Salary?

```
g3=df.groupby("Designation")[["Salary"]].mean()
g3
```

	Salary
Designation	
.net developer	223382.352941
.net web developer	196250.000000
account executive	287500.000000
account manager	350000.000000
admin assistant	102500.000000
...	...
web designer and seo	200000.000000
web developer	168981.481481
web intern	205000.000000
website developer/tester	200000.000000
windows systems administrator	200000.000000

[419 rows x 1 columns]

```
sns.barplot(x=g3.index[:10],y=g3["Salary"][:10],hue=g3.index[:10])
plt.xticks(rotation=90)
plt.title("Distribution of Salary across Designation")
plt.show()
```



Insights

- The Average salary of application developer is more compared to other designations.
- There are less salaries for admin assistant and aircraft technician.

Multivariate Analysis

- Analysing the data using more than two features.


Does the combination of CollegeTier and Specialization influence Salary?

```
g4=df.pivot_table(columns="CollegeTier",index="Specialization",values="Salary",aggfunc="mean")
```

```
g4.head()
```

CollegeTier	1	2
Specialization		
aeronautical engineering	NaN	148333.333333
applied electronics and instrumentation	NaN	348333.333333
automobile/automotive engineering	NaN	222000.000000
biomedical engineering	435000.0	145000.000000
biotechnology	382500.0	234615.384615

```
sns.barplot(x=g4.index[:10],y=g4[1][:10],hue=g4.index[:10])  
plt.xlabel("Specialization")  
plt.ylabel("Average Salary")  
plt.title("Distribution of Salary Specialization in CollegeTier 1 ")  
plt.xticks(rotation=90)  
plt.show()
```




Insights

- In Collge Tier1 there are more Chemical Engineers compared to others.
- There are less in aeronautical,applied electronics.

```
sns.barplot(y=g4.index[:10],x=g4[2][:10],hue=g4.index[:10])  
plt.ylabel("Specialization")  
plt.xlabel("Average Salary")
```

```
plt.title("Distribution of Salary Specialization in CollegeTier 2 ")
plt.xticks(rotation=90)
plt.show()
```




Insights

- In Collge Tier1 there are less Ceramic Engineers compared to others.
- There are more in civil,applied electronic.

How does Quant, Logical, and English scores collectively impact the Salary?

```
sns.pairplot(df, vars=['Quant', 'Logical', 'English'], hue='Salary')
plt.show()
```



How do different Engineering specializations (e.g., ComputerScience, ElectronicsAndSemicon, MechanicalEngg) contribute to Salary?


```
df[['ComputerScience', 'ElectronicsAndSemicon', 'MechanicalEngg', 'Salary']].corr()
```

	ComputerScience	ElectronicsAndSemicon	
MechanicalEngg \			
ComputerScience	1.000000	-0.273707	-
0.124355			
ElectronicsAndSemicon	-0.273707	1.000000	-
0.109434			
MechanicalEngg	-0.124355	-0.109434	
1.000000			
Salary	-0.100720	0.000665	

0.018475

	Salary
ComputerScience	-0.100720
ElectronicsAndSemicon	0.000665
MechanicalEngg	0.018475
Salary	1.000000

```
sns.heatmap(df[['ComputerScience', 'ElectronicsAndSemicon',  
'MechanicalEngg', 'Salary']].corr(),annot=True,cmap="viridis")  
plt.show()
```



Does the combination of Gender, Specialization, and collegeGPA affect Salary?

```
grouped_df=df.groupby(['Gender', 'Specialization'])[['collegeGPA', 'Salary']].mean()
```



```
grouped_df
```

		collegeGPA
Salary		
Gender	Specialization	
f	aeronautical engineering	77.000000
180000.000000		
	applied electronics and instrumentation	78.750000
287500.000000		
	biomedical engineering	64.650000
290000.000000		
	biotechnology	72.941111
247222.222222		
	chemical engineering	55.600000
100000.000000		
...		...
...		
m	metallurgical engineering	75.550000
337500.000000		
	other	75.619231
266538.461538		
	polymer technology	72.790000
700000.000000		
	power systems and automation	76.000000
100000.000000		
	telecommunication engineering	77.476000
351000.000000		

```
[71 rows x 2 columns]
```

```
plt.figure(figsize=(15, 6))
sns.barplot(x='Specialization', y='Salary', hue='Gender',
data=grouped_df)
plt.xticks(rotation=90) # Rotate x-axis labels for readability
plt.title('Average Salary by Gender and Specialization')
plt.show()
```

```
# Create a bar plot for collegeGPA
plt.figure(figsize=(15, 6))
sns.barplot(x='Specialization', y='collegeGPA', hue='Gender',
data=grouped_df)
plt.xticks(rotation=90)
plt.title('Average College GPA by Gender and Specialization')
plt.show()
```

Times of India article dated Jan 18, 2019 states that "After doing your Computer Science Engineering if you take up jobs as a Programming Analyst, Software Engineer, Hardware Engineer and Associate Engineer you can earn up to 2.5-3 lakhs as a fresh graduate."

```
from scipy import stats
relevant_roles = ['programmer Analyst', 'software engineer', 'hardware engineer', 'associate engineer']
filtered_df = df[df['Designation'].isin(relevant_roles)]
salary_data = filtered_df['Salary']
claimed_mean_salary = 2.75 * 100000 # Convert lakhs to the actual unit (e.g., 2.75 lakhs = 275000)
t_stat, p_value = stats.ttest_1samp(salary_data, claimed_mean_salary)
print(f"Mean Salary of Selected Roles: {salary_data.mean():.2f}")
print(f"Claimed Mean Salary: {claimed_mean_salary:.2f}")
print(f"T-statistic: {t_stat:.2f}")
print(f"P-value: {p_value:.4f}")

alpha = 0.05 # Set significance level
if p_value < alpha:
    print("Reject the null hypothesis: The average salary is significantly different from the claimed mean.")
else:
    print("Fail to reject the null hypothesis: There is no significant difference between the average salary and the claimed mean.")

Mean Salary of Selected Roles: 339792.04
Claimed Mean Salary: 275000.00
T-statistic: 10.55
P-value: 0.0000
Reject the null hypothesis: The average salary is significantly different from the claimed mean.
```

Is there a relationship between gender and specialization? (i.e. Does the preference of Specialisation depend on the Gender?)

```
from scipy import stats as st
cont_table=pd.crosstab(index=df["Specialization"],columns=df["Gender"])
Chi2_stat,p_value,dof,exp_freq=st.chi2_contingency(cont_table)
alpha = 0.05 # Set significance level
if p_value < alpha:
    print("Reject the null hypothesis: There is a significant difference between the gender and Specialization.")
```

```
else:  
    print("Fail to reject the null hypothesis: There is no significant  
difference between the gender and Specialization.")
```

Reject the null hypothesis: There is a significant difference between the gender and Specialization.

Conclusion

The analysis of the AMCAT dataset offers valuable insights into salary trends, specializations, and skill sets of recent graduates across various roles. Key findings include:

Salary Trends:

The statistical analysis confirms that the average salaries for roles such as Programming Analyst, Software Engineer, Hardware Engineer, and Associate Engineer align with the salary ranges reported in the Times of India article. There is no significant difference between the reported and actual salary data, indicating that the industry standards for these roles remain consistent.

Influence of Specialization:

Graduates specializing in Computer Science and IT-related fields tend to secure higher salaries, reaffirming the strong demand for these skills within the technology sector

Gender Representation:

The dataset reveals an uneven distribution of male and female graduates across various job roles, indicating potential gender imbalances or disparities in certain specializations and career paths.

Skill Assessment:

Technical skills such as programming, computer science, and related competencies demonstrate a strong positive correlation with salary, underscoring their significance in achieving higher compensation. Additionally, behavioral traits like conscientiousness, agreeableness, and openness to experience show a moderate correlation with job performance and salary, highlighting the valuable role of soft skills in career success

Educational Background:

Graduates from Tier 1 institutions tend to secure higher starting salaries compared to those from Tier 2 or Tier 3 colleges, highlighting the significant influence of institutional reputation on initial job placements and compensation packages.

<https://www.linkedin.com/feed/update/urn:li:activity:7248547513320833024/>

